

AI Voice Generation

Turning text into speech — and using it responsibly

What You Will Learn

- Understand the difference between text-to-speech and voice cloning.
- Know where businesses actually use AI voice (IVR, e-learning, ads, dubbing, accessibility).
- Match voice tone and pace to the right audience.
- Apply the consent and disclosure rules before using AI voice commercially.

1. What Is AI Voice Generation?

AI voice generation turns written text into spoken audio using a computer-generated voice. Type a script, pick a voice, and the tool reads it aloud in seconds — no microphone, no recording studio, no human voice actor needed. There are two related but different things people call "AI voice":

- **Text-to-Speech (TTS):** you type text, AI speaks it in a pre-built voice (e.g. "Aria", "Rohan").
- **Voice cloning:** AI learns the sound of a specific real person's voice (from a sample recording) and can then make that exact voice say new things it never actually said.

Example — typing a script

You type: "Welcome to CampusBrew — your favourite study spot, now with faster Wi-Fi." A TTS tool reads this aloud in a warm, friendly voice in about 5 seconds, ready to drop into an Instagram Reel.

2. Where Businesses Actually Use It

- **IVR / phone systems:** "Press 1 for billing, press 2 for support" — recorded once, updated instantly by editing text instead of re-recording.
- **E-learning narration:** turning a course PDF or slide deck into a narrated video without hiring a narrator.
- **Podcast and ad voiceovers:** quick draft voiceovers for testing an ad script before paying for a professional recording.
- **Multilingual dubbing:** the same video narrated in Hindi, English, and Tamil without three separate recording sessions.
- **Accessibility:** reading website content aloud for visually impaired users.

3. Picking the Right Voice for the Message

The same script sounds completely different depending on the voice's tone, pace, and accent. A voice that works for one audience can feel wrong for another — this is a choice, not a default setting to ignore.

Example — same script, wrong voice

A college coffee-shop ad script read in a slow, formal, corporate-sounding voice will feel out of place for an 18-24 year old audience — even though every word is identical to a version read in an upbeat, casual voice. The words didn't change; the fit did.

4. The Ethics Line: Consent and Disclosure

Voice cloning raises a real risk: someone's voice can be recreated to say things they never said. Two rules matter for business use:

- 1 **Consent:** never clone a real person's voice (a colleague, a customer, a public figure) without their explicit permission.
- 2 **Disclosure:** if a voice in an ad, video, or customer-facing message is AI-generated, say so — presenting it as a real human recording when it isn't is misleading.

ACTIVITY 1: Match the Voice to the Audience

Step 1. Write a 30-second script for a business of your choice (e.g. CampusBrew, a gym, a local bakery).

Step 2. Generate it using two very different voice settings — one formal/corporate, one casual/friendly.

Step 3. Listen to both and decide which one actually fits your target audience, and why.

Worked example: Script: "Fresh coffee, faster Wi-Fi, and a spot to actually get work done. That's CampusBrew." Read in a formal, deep, slow voice it sounds like a bank ad. Read in an upbeat, quick, casual voice it sounds like it's actually talking to a student. Same words, two very different fits.

ACTIVITY 2: Fix a Robotic-Sounding Script

Step 1. Write a sentence with no punctuation cues, e.g. "our coffee is fresh and our staff are friendly and we are open till 10pm".

Step 2. Generate the audio and notice how flat and run-on it sounds.

Step 3. Rewrite it with commas, shorter sentences, and an exclamation point where it matters, then regenerate.

Worked example: Before: "our coffee is fresh and our staff are friendly and we are open till 10pm" (reads as one flat run-on). After: "Our coffee is fresh. Our staff are friendly. And we're open till 10pm!" (reads with natural pauses and energy in the right place).

ACTIVITY 3: Draft an AI-Voice Disclosure

Step 1. Imagine you're using an AI voice in a customer-facing video ad.

Step 2. Write one honest, short sentence disclosing that the voice is AI-generated.

Step 3. Identify one scenario where you would need a real person's explicit written consent before cloning their voice for business use.

Worked example: *Disclosure line: "This ad features an AI-generated voice." Consent scenario: cloning your company's founder's voice for a recruiting video — even though it's your own company, it's still that specific person's voice, so you need their sign-off before using an AI clone of it.*