

AI for HR

Faster drafts, human judgment on bias and accuracy

What You Will Learn

- Know where AI genuinely speeds up HR work.
- Understand why AI resume screening can inherit bias from past hiring data.
- Write a job description that avoids exclusionary language and inflated requirements.
- Fact-check AI-drafted HR content against the real policy it's based on.

1. Where AI Actually Helps HR

- **Job description drafting:** turning a rough bullet list of duties into a polished, complete JD.
- **Policy explainer drafts:** a first-pass plain-language explanation of a leave policy, expense policy, etc.
- **First-pass resume screening:** quickly checking resumes against a role's stated requirements.
- **HR FAQ content:** drafting answers to common employee questions for an internal handbook.

2. The Bias Risk in AI Resume Screening

AI resume-screening tools are often trained on a company's past hiring data. If that past data reflects biased hiring patterns (favoring certain schools, name patterns, or backgrounds), the AI can learn to repeat that bias — consistently, at scale, without anyone deciding it should. This is why AI resume screening must be audited by a human, not trusted as a neutral, objective filter.

Example — biased screening

An AI tool trained mostly on resumes of past hires from a handful of well-known universities may consistently rank equally qualified candidates from lesser-known colleges lower — not because of any actual skill difference, but because the training data taught it that pattern. Left unchecked, this repeats and hardens a bias nobody explicitly chose.

3. Writing an Inclusive Job Description

Certain words and phrases in job postings unintentionally discourage some qualified candidates from applying. AI can help draft a JD quickly — but it can also repeat these patterns unless you specifically ask it to avoid them.

- Avoid gendered or exclusionary language: "rockstar", "ninja", "young and energetic team".
- Avoid inflated requirements that filter out good candidates for no real reason: "5+ years experience" for an entry-level role.
- Be specific about actual must-have skills vs nice-to-have skills, instead of one long undifferentiated list.

ACTIVITY 1: Draft and Audit a Job Description

Step 1. Write a rough bullet list of duties for a role (e.g. Marketing Intern).

Step 2. Ask AI to turn it into a full job description.

Step 3. Read it back specifically looking for biased or exclusionary language, and fix anything you find.

Worked example: Rough bullets: "manage social media, help with events, some Excel work". AI draft might include a phrase like "looking for a young, energetic self-starter" — flag "young" as age-exclusionary language and replace with "motivated self-starter", which describes the same trait without excluding older candidates.

ACTIVITY 2: Separate Must-Haves From Nice-to-Haves

Step 1. Ask AI to list 3 must-have and 3 nice-to-have skills for a role of your choice.

Step 2. Review the must-have list critically: is each one actually necessary, or is it an inflated requirement?

Step 3. Rewrite any must-have that's actually artificially strict (e.g. years of experience for an entry-level role) as a nice-to-have instead.

Worked example: Role: Junior Data Analyst. AI's must-haves might include "3+ years of experience with SQL" — for a junior role, that is an inflated requirement. Move it to nice-to-have, and replace the must-have with "basic understanding of SQL, willingness to learn on the job".

ACTIVITY 3: Fact-Check an AI-Drafted Policy FAQ

Step 1. Give AI a real (or sample) company leave/expense policy and ask it to draft one FAQ answer from it.

Step 2. Compare the AI answer line-by-line against the actual policy text.

Step 3. Flag anything AI added that wasn't actually stated in the policy you gave it.

Worked example: Policy states: "Employees get 12 paid leave days per year." If AI's FAQ answer adds "...and unused days carry over to the next year" without that ever being stated in the policy, that's an invented detail that must be removed or verified with actual HR before it's published.